Eric Salerno — Full Career Portfolio

A comprehensive look at my career, for anyone who wants more than a resume — recruiters without a specific job description, hiring managers doing early diligence, or anyone curious what "technical PM who builds" actually means in practice. Skip to whatever's relevant, or read straight through — the homepage has the short version, and the AI project case studies are written up in full there too.

About

I'm a technical product manager who builds. I don't just own the backlog — I design the architecture and ship the system. Most recently, at Amperity, I designed and built a two-part AI-assisted integration system — a HITL discovery pipeline and a dockerized execution engine — that automates third-party API integration and cut development time roughly 90%, from about four weeks down to a couple of days.

Over nearly 30 years I've gone from senior full-stack developer to technical product/program manager to Director and Senior Director of Technical Program Management, and the through-line has stayed the same: turn an idea — or a model, or a spec — into a shipped, integrated product. I shipped the PlayReady DRM platform APIs at Microsoft, built fraud-and-abuse systems for Amazon Coins, led a data-exchange and AI-NLP partnership with Epic at Press Ganey, and spent the last year owning a connector and activation program at Amperity — shipping 24 connectors and the agentic tooling around them.

What I care about now: agentic workflows, LLM orchestration, and autonomous systems that do real work in production — built fast, on small teams, where the PM is close enough to the code to debate the architecture and close enough to the customer to know it matters.

Amperity

Technical Product Manager · Feb 2025 – Jun 2026 · Remote / US — Enterprise Customer Data Platform (CDP) & Data Activation

Amperity is an enterprise customer-data-platform and data-activation company. I owned the connector and destination-activation program end-to-end — from customer demand intake through requirements, spec authoring, engineering hand-off, release tracking, and adoption reporting — across a ~16-month tenure (Feb 2025 – Jun 2026).

Owning the Connector & Destination Integration Program

The company's value depends on activating customer data into a long tail of third-party destinations — ad platforms, ESPs, marketing and loyalty tools — but demand arrived from every direction (sales, customer success, customers directly) with no consistent way to capture, judge, scope, or report on it. I stood up and ran the end-to-end intake pipeline: captured demand through the product feedback tool and direct customer signals, translated it into structured, engineering-ready requests, and prioritized against business context. I authored 59+ connector request specs spanning CAPI/server-side event APIs, ESPs, ad networks, messaging, loyalty, and real-time streaming destinations, enriched each with the business context engineering needed to prioritize — which customers wanted it, urgency, strategic fit — and maintained the release roadmap (real-time / traditional / pilot swimlanes) presented at the recurring product-monthly meeting. The result: a chaotic, multi-channel demand stream became a single governed pipeline with clear status, ownership, and a forward-looking roadmap — the connective tissue between customers, sales/CS, and engineering. Over the tenure this shipped 24 connectors to production (12 batch, 12 real-time).

Writing Specs Engineering Could Build From

Engineering needed crisp, validated specs to build connectors efficiently — without them, build cycles stalled on ambiguity and rework. I wrote the full ladder of product artifacts — business requirements documents, functional specs, and engineering hand-off specifications — for destination connectors and platform capabilities, and standardized them into repeatable templates covering problem framing, MUST/SHOULD/MAY requirements, scope and out-of-scope, dependencies, risks, and open questions, so every spec hit the same bar. Before committing engineering time, I validated integration feasibility myself: interrogating partner API docs, building Postman collections, and standing up mock APIs to prove a spec was buildable before anyone wrote production code. I also authored platform-level requirements for several bigger bets — a self-serve connector capability, a connector SDK, real-time connectors, and a destination testing/"revamp" initiative. The result was a pipeline of specs that arrived scoped, de-risked, and validated, letting engineering take on a much wider catalog of integrations than they otherwise could.

Building an AI-Assisted Integration System — and the Tooling Around It

Program ownership at this scale generated enormous amounts of repeatable analytical and reporting work, and the most technical response I built to it was a two-part AI-assisted integration system in Python and Claude. The discovery half is a human-in-the-loop pipeline: it takes an OpenAPI spec and sample data, uses web search to understand the destination API, validates the auth model with a live credential ping, and produces a confidence-scored mapping from source columns to destination endpoints — which I could review and adjust — that gets written out as a config file fully specifying the integration. The execution half is a dockerized engine that exposes an endpoint, loads that config, applies the required transformations, and delivers the data to the destination, handling production concerns like rate limiting, payload size limits, callbacks, and pagination, backed by a custom-built credential vault and a SQL database for config versioning and run history. I held to one architecture principle throughout: deterministic Python everywhere the logic was unambiguous, and LLM calls only where natural-language reasoning actually added value. I validated the resulting config format against Amperity's existing production integration library and found it 85% compatible out of the box, with the remaining 15% blocked by unsupported auth mechanisms or SOAP-based APIs. Along the way I evaluated and discarded both a RAG layer and a multi-agent orchestration approach — neither solved the actual problem, and a single focused pipeline did the job without the added complexity.

Beyond that system, I built more than 30 reusable automation tools and workflows covering release status reporting, requirements enrichment, blocked-item triage, documentation auditing, roadmap generation, and analytical-narrative generation, plus internal MCP servers that connected the team's tools into automated workflows, and six mock destination APIs built with FastAPI so integrations to credential-gated partners could be scoped and demoed without needing live access.

Turning Program Data into Decisions

Product leadership, customer success, and customers all needed visibility into integration adoption, activation and usage, and program health, but the data was scattered across systems. I designed and shipped multiple production dashboards and data apps — built in Streamlit-in-Snowflake and in HTML — covering connector release status, activation and usage trends, monthly-business-review metrics, and a documentation/knowledge dashboard for the team. I applied real analytical rigor to each: choosing the visual encoding that actually fit the question, and using statistical and anomaly-detection thinking rather than naive fixed thresholds for monitoring. I calibrated every output to its audience, building executive-ready roadmap decks and MBR references alongside more operational, technical views for the team itself.

Closing the Loop Between Customers and Engineering

Customers and field teams regularly hit friction operating integrations and activation features, and sales needed fast, credible feasibility answers. I triaged reported connector and capability gaps by verifying the claim across source code, vendor docs, product docs, live tenant configuration, and run history — distinguishing a genuine product gap from a misconfiguration or a documentation gap before anyone escalated it as a bug. I built journey and campaign automation tooling and validated activation mechanics like templating, segmentation, and scheduling against real and synthetic data, and produced demo and pilot environments and synthetic datasets to prove out features and integrations for prospects and customers. That work shortened the loop between a customer reporting a problem and knowing exactly whether it was a config fix, a docs fix, or a build — and gave the field answers it could trust.

Running the Operational Backbone

The integration program touched engineering, product, sales, customer success, docs, and partners, and I was the hub connecting all of them — gathering demand, aligning priorities, tracking delivery, and reporting status on recurring cadences that included weekly status updates, a monthly roadmap review, and MBRs. I managed a tracking-system migration mid-program, consolidating the team onto a new project/board structure and re-pointing all the automation that depended on the old one. I defined a repeatable process running from intake through spec authoring, engineering hand-off, release, and adoption reporting, and encoded it as tooling rather than leaving it as tribal knowledge, so it kept working as the team and the program grew.

Across the tenure, the headline numbers: the AI-assisted integration system cut connector development time roughly 90%, from about four weeks down to about two days; its output was 85% compatible with Amperity's existing production integration library out of the box; 24 connectors shipped to production (12 batch, 12 real-time); intake automation cut the time from a customer request to a dev-ready spec and ticket from about a week to about an hour; and three hand-built pilot connectors kept Amperity competitive in active sales deals.

Press Ganey

Senior Director, Technical Program Management · Mar 2022 – Aug 2024 · Greater Seattle — promoted from Distinguished TPM; reported directly to the CTO

Press Ganey is a leading healthcare experience and performance improvement company serving thousands of healthcare organizations. I joined as a Distinguished TPM and was promoted to Senior Director, owning a portfolio of strategic initiatives directly for the CTO.

Leading the Epic Partnership & Integration

Press Ganey's patient-experience data had significant untapped value inside the Epic EHR ecosystem — the dominant system in US healthcare — but no integration existed, and building one required a true strategic partnership across two organizations with different cultures, timelines, and architectures. I spearheaded the Press Ganey/Epic integration, delivering key milestones including MyChart (Epic's patient-facing portal) and Cheers CRM (Epic's relationship-management tool), bringing Press Ganey data directly into the systems clinicians and administrators actually use. I was the primary relationship owner across both organizations — managing senior leadership and key decision-makers at both companies simultaneously, coordinating product, engineering, sales, and marketing at both to keep execution aligned with the partnership roadmap, gathering and analyzing market requirements to define scope, and presenting executive summaries and status reports to senior leadership at both companies. The result was a live Epic integration giving healthcare organizations real-time patient-experience insights inside the clinical workflow — a differentiated capability that expanded Press Ganey's EHR footprint.

Unifying Technology Across Acquisitions

Press Ganey had made multiple acquisitions, each with its own tech stack, data architecture, and engineering team, all needing assessment, rationalization, and unification onto the Press Ganey platform without disrupting live products. I drove the technical unification across the acquired companies, working directly with the CTO on integration strategy. I led technical due diligence for multiple acquisitions — assessing feasibility, integration risk, and synergy opportunities before commitments were made — developed and implemented migration strategies for acquired systems and data, coordinating engineering teams across multiple locations, and built and presented technical roadmaps and integration proposals to senior leadership and board members.

Building the TPM Career Path from Scratch

Press Ganey had no defined TPM discipline — no career path, no role definitions, no performance expectations — which made it impossible to hire, develop, or retain technical program management talent consistently. I built the first TPM career path at Press Ganey from scratch, working closely with the CTO and their direct reports to define roles, responsibilities, and career progression for ICs and managers. I authored job descriptions, role definitions, and career progression models for every level, defined performance expectations, KPIs, and compensation guidelines, and delivered presentations and workshops to build organization-wide buy-in, including transitioning existing Business Analysts and Project Managers into TPM roles. The result was a structured TPM discipline across the engineering organization, and the foundation for consistent hiring and development going forward.

Making the Call Not to Pursue FedRAMP

Press Ganey was considering FedRAMP authorization to expand into the federal healthcare market — a substantial, typically multi-year, multi-million-dollar investment that needed rigorous analysis before committing. I led the feasibility analysis across technical, operational, and financial requirements, evaluated the full compliance scope — infrastructure, process controls, staffing, tooling, third-party assessment costs, and ongoing maintenance — and built an ROI-based no-go recommendation that I drove through to an actual decision rather than a report that sat on a shelf. Leadership accepted the no-go recommendation, avoiding a costly multi-year program that wouldn’t have generated sufficient return — one of the harder and more valuable calls a TPM can make and defend.

Unifying the SDLC Around Jira and Jira Align

Engineering teams operated with inconsistent processes and tooling, making it impossible to get reliable visibility into project health, team performance, or roadmap progress at the executive level. I drove adoption of Jira and Jira Align across multiple engineering teams, collaborated with executive management and development leadership to define and prioritize KPIs, built custom reports and dashboards tracking progress, velocity, and milestones, and presented data-driven insights to leadership on a recurring basis.

Accolade

Director of TPM / Technical Product Manager · Mar 2017 – May 2022 · Seattle, WA — promoted from IC TPM to Director, Mar 2021

Accolade is a personalized health-navigation company serving large employer clients. As an IC TPM I led the design and implementation of new data-ingestion infrastructure for partner feeds — member claims, employment data, enrollments — establishing the data-quality foundation later projects depended on. I then owned a multi-year migration from a heavily customized Microsoft Dynamics CRM to a custom in-house solution for Accolade's 50+ Health Assistant workforce, using a phased rollout (pilot → structured feedback → iterate → full scale) that improved user satisfaction and data accuracy without disrupting Health Assistants serving members day to day. When Accolade acquired 2nd.MD and PlushCare, I drove the program to unify Accolade's core technology and infrastructure with both acquired companies — data mapping, API design, HIPAA-compliant data security, and integration across services/data tiers, client portals, CRM portals, and member-facing portals. I was promoted to Director of TPM and led a team of 4-5 Technical Program Managers responsible for all data ingress and egress across the platform, building deep hands-on expertise across data warehouses, data lakes, CRM systems, and partner feeds to resolve client and partner data issues quickly.

Amazon

Senior Technical Program Manager · Oct 2014 – Jul 2015 · Seattle, WA — Amazon Coins team

Amazon Coins is Amazon's virtual currency for purchasing apps, games, and in-app content. I developed and implemented fraud-detection models using data-mining techniques to catch abuse patterns specific to virtual currency — patterns traditional retail fraud models weren't built to catch — using a customer-experience-first approach that improved detection coverage without adding friction for legitimate customers. I also built customer-segmentation models to identify high-value customers ("whales") within the user base, enabling marketing to target campaigns calibrated to actual customer behavior and value.

Microsoft

Senior Program Manager · Jan 2004 – Oct 2014 (two stints, with Appature in between Nov 2012 – Sep 2013) · Redmond, WA

I led the initial API design and shipped the PlayReady DRM platform — core functionality and APIs for Windows 8 and Silverlight 2/3/4 — a true zero-to-one platform build that needed Tier-1 content-partner adoption against a contested standards landscape. I drove product strategy end-to-end (scenarios, feature specs, Sprint-based Agile delivery), fostered collaboration across Windows, IIS, A/V Services, and Online Services, engaged directly with Tier-1 and third-party partners, and represented PlayReady at industry standards bodies. Later, on Xbox One, I led the security program hardening the console's offline licensing model at the silicon and firmware level — security updates to the optical disc drive DSP, coordinated across ODD teams, silicon vendors, and firmware engineers, plus third-party penetration testing before launch — and led the launch of a security-focused SaaS delivery model for critical Microsoft technologies. Earlier in my Microsoft tenure, I was product owner and PM for a highly customized SharePoint-based knowledge-sharing platform for Microsoft Consulting Services, leading a 20+ person offshore development and testing team in Bangalore and traveling to India twice for major delivery milestones.

Earlier Roles

Appature — Technical Program Manager, Nov 2012 – Sep 2013, Seattle, WA (martech startup, acquired by IMS Health). Owned the platform APIs for a Marketing Relationship Management SaaS on AWS, used by customers, internal developers, and Customer Success; drove integration of the platform's ETL/CRUD APIs into SnapLogic's visual designer to reduce customer adoption friction.

Concur Technologies — Program Manager. Grew from individual-contributor developer to dev lead to PM; supported Fortune-500 customer engagements.

EDS / General Motors — Product Manager / Developer / Test. Early ASP web-app adopter.

AI Portfolio Projects

Since mid-2026 I've been building and directing AI coding agents on a portfolio of from-scratch technical projects — full write-ups and code are linked below; this is the short version.

Agent-Reviewed Mock Interview Tool

A live interview-prep tool built with Claude Code: real answers saved through the browser, reviewed and reacted to by the agent in real time — draft/approve/feedback loop, live status updates, and a self-caught reviewer bias that got fixed and turned into a reusable Claude Code skill.

Case study →

AI-Assisted Martech Integration Builder

An agent ingests an OpenAPI spec, cross-references domain knowledge files for gotchas the spec doesn't mention, gates risky operations behind human review, then generates and verifies a working MCP tool server against a live account (Klaviyo and HubSpot).

GitHub → · Live demo →

RAG Pipeline, Built From Scratch

A retrieval-augmented generation system with no LangChain, no LlamaIndex, no vector database. A custom parser chunks OpenAPI specs at the operation level so schema and meaning stay together — every stage, from chunking to prompt assembly, is visible and attributable to project code.

GitHub →

judge-dread

A local-first regression-testing harness for AI prompts: deterministic checks plus LLM-as-judge evaluation, built entirely on Gemini's free tier.

GitHub →

nanoGPT, Trained From Scratch

A character-level GPT trained from scratch on public-domain science fiction from Project Gutenberg, built with nanoGPT — hands-on with the fundamentals of how these models actually learn.

GitHub →

learnpath-agent

An adaptive learning-path agent modeled on Amazon Ads Academy's real catalog taxonomy: an LLM plans and re-plans a personalized course sequence through a 54-item AI-concepts catalog based on a learner's goal and quiz scores, showing exactly which candidates it considered and why — with a deterministic fallback if the LLM call fails.

Case study → · GitHub →

Raspberry Pi Home Audio Platform

One control surface for a six-zone home audio system — two Marantz AV receivers and an AudioControl multi-zone amp — running on a Raspberry Pi. I reverse-engineered the amplifier's undocumented HTTP/JSON control API, wrote the Python device layer and a FastAPI web UI, and wired a Gemini "systems investigator" agent into the repo: it answers a pull-request comment that tags it ("@gemini") with a diagnosis and a proposed fix, with three-stage model failover to ride out rate limits.

GitHub →

Skills & Credentials

Product / Program craft: requirements (BRD/PRD/FSD), product discovery and customer interviews, prioritization, roadmapping, release management, stakeholder management, executive communication, process design, MBR/QBR reporting.

Technical: Python, SQL (Snowflake), REST API integration and analysis, FastAPI, Postman/Newman, data visualization (Streamlit, HTML/D3), CI/CD pipelines, Git, LLM/AI tooling (MCP servers, prompt/skill authoring), mock-service and test-harness design, data modeling for analytics.

Domain: Customer Data Platforms (CDP), marketing/advertising activation, data integration and connectors, identity/data onboarding, ad-tech and martech ecosystems (CAPI/server-side events, ESPs, ad networks, loyalty, messaging), healthcare data interoperability (Epic).

AI / ML (agent-directed, from-scratch, ongoing portfolio): transformer and self-attention architecture, tokenization design, training loops and optimization, PyTorch, reproducible ML pipeline engineering, RAG architecture from first principles, LLM-as-judge evaluation design, structured-output API usage, regression-testing harness design, multi-agent development workflows, agentic tool design (MCP server generation, HITL workflows).

Education: University of Michigan — B.S., Industrial & Operations Engineering. Michigan Technological University — Engineering.

Certifications: AI/ML Foundations (Neural Networks, Machine Learning, Thinking Machines) · Recommendation Systems with Python · Predictive Analytics & Data Mining.

Resume & Contact

Resume: Download Resume (PDF)

Contact: esalerno86@gmail.com · LinkedIn · GitHub